

Hackathon II

The Evolution of Todo – Mastering Spec-Driven Development & Cloud Native AI

The future of software development is **AI-native and spec-driven**. As AI agents like Claude Code become more powerful, the role of the engineer shifts from "syntax writer" to "system architect." We have already explored Spec-Driven Book Authoring. Now, we want you to master the **Architecture of Intelligence**.

In this hackathon, you will master the art of building applications iteratively—starting from a simple console app and evolving it into a fully-featured, cloud-native AI chatbot deployed on Kubernetes. This journey will teach you the [Nine Pillars of AI-Driven Development](#), [Claude Code](#), [Spec-Driven Development with Reusable Intelligence](#) and Cloud-Native AI technologies through hands-on implementation.

Excel in the Hackathon and Launch Your Journey as an AI Startup Founder

We've recently launched **Panaversity** (panaversity.org), an initiative focused on teaching cutting-edge AI courses. If you perform well in this hackathon, you may be invited for an interview to join the **Panaversity core team** and potentially step into the role of a **startup founder** within this growing ecosystem. You will get a chance to work with Panaversity founders Zia, Rehan, Junaid, and Wania and become the very best. You may also get a chance to teach at Panaversity, PIAIC, and GIAIC.

What You Will Learn

- Spec-Driven Development using Claude Code and Spec-Kit Plus
- Reusable Intelligence: Agents Skills and Subagent Development
- Full-Stack Development with Next.js, FastAPI, SQLModel, and Neon Serverless Database
- AI Agent Development using OpenAI Agents SDK and Official MCP SDK
- Cloud-Native Deployment with Docker, Kubernetes, Minikube, and Helm Charts
- Event-Driven Architecture using Kafka and Dapr
- AIOps with kubectl-ai, kagent and Claude Code
- Develop Cloud-Native Blueprints for Spec-Driven Deployment

Research Note: Deployment Blueprints for Spec-Driven Deployment

1. [Is Spec-Driven Development Key for Infrastructure Automation?](#)
2. [ChatGPT Progressive Learning Conversation](#)
3. [Spec-Driven Cloud-Native Architecture: Governing AI Agents for Managed Services with Claude Code and SpecKit](#)

Requirements

You are required to complete the **5-Phase "Evolution of Todo" Project** using Claude Code and Spec-Kit Plus. The core deliverables are:

- **Spec-Driven Implementation:** You must implement all **5 Phases** of the project (detailed below). You are strictly required to use **Spec-Driven Development**. You

must write a Markdown Constitution and Spec for every feature of the phase, and use **Claude Code** to generate the implementation.

Constraint: You cannot write the code manually. You must refine the Spec until Claude Code generates the correct output.

- **Integrated AI Chatbot:** In Phases III, IV, and V, you must implement a conversational interface using **OpenAI Chatkit**, **OpenAI Agents SDK**, and **Official MCP SDK**. The bot must be able to manage the user's Todo list via natural language (e.g., "Reschedule my morning meetings to 2 PM").
- **Cloud Native Deployment:** In Phases IV and V, you must deploy the chatbot locally on Minikube, and on the cloud on DigitalOcean Kubernetes (DOKS).

Todo App Feature Progression

Basic Level (Core Essentials)

These form the foundation—quick to build, essential for any MVP:

1. Add Task – Create new todo items
2. Delete Task – Remove tasks from the list
3. Update Task – Modify existing task details
4. View Task List – Display all tasks
5. Mark as Complete – Toggle task completion status

Intermediate Level (Organization & Usability)

Add these to make the app feel polished and practical:

- 1.
2. Priorities & Tags/Categories – Assign levels (high/medium/low) or labels (work/home)
3. Search & Filter – Search by keyword; filter by status, priority, or date
4. Sort Tasks – Reorder by due date, priority, or alphabetically

Advanced Level (Intelligent Features)

1. Recurring Tasks – Auto-reschedule repeating tasks (e.g., "weekly meeting")
2. Due Dates & Time Reminders – Set deadlines with date/time pickers; browser notifications

Hackathon Phases Overview

Phase	Description	Technology Stack	Points	Due Date
Phase I	In-Memory Python Console App	Python, Claude Code, Spec-Kit Plus	100	Dec 7, 2025
Phase II	Full-Stack Web Application	Next.js, FastAPI, SQLModel, Neon DB	150	Dec 14, 2025
Phase III	AI-Powered Todo Chatbot	OpenAI ChatKit, Agents SDK, Official MCP SDK	200	Dec 21, 2025
Phase IV	Local Kubernetes Deployment	Docker, Minikube, Helm, kubectl-ai, kagent	250	Jan 4, 2026
Phase V	Advanced Cloud Deployment	Kafka, Dapr, DigitalOcean DOKS	300	Jan 18, 2026
TOTAL			1,000	

Bonus Points

Participants can earn additional bonus points for exceptional implementations:

Bonus Feature	Points
Reusable Intelligence – Create and use reusable intelligence via Claude Code Subagents and Agent Skills	+200
Create and use Cloud-Native Blueprints via Agent Skills	+200
Multi-language Support – Support Urdu in chatbot	+100
Voice Commands – Add voice input for todo commands	+200
TOTAL BONUS	+600

Timeline

- Submission Deadline:** On Sundays on dates as mentioned above.
- Live Presentations:** On Sundays, December 7, 14, and 21, 2025 and on January 4 and 18, 2026 starting at 8:00 PM on Zoom. Final Live Presentation date to be determined.

Top submissions will be invited via WhatsApp to present live on Zoom.

Note: All submissions will be evaluated. Live presentation is by invitation only, but does not affect final scoring.

Milestone	Date	Description
Hackathon Start	Monday, Dec 1, 2025	Documentation released
Phase I Due	Sunday, Dec 7, 2025	Console app checkpoint
Phase II Due	Sunday, Dec 14, 2025	Web app checkpoint
Phase III Due	Sunday, Dec 21, 2025	Chatbot checkpoint
Phase IV Due	Sunday, Jan 4, 2026	Local K8s checkpoint
Final Submission	Sunday, Jan 18, 2026	All phases complete
Live Presentations	Sundays, Dec 7, 14, 21, and Jan 4 and 18	Top submissions present

Submit and Present Your Project:

Once you have completed the project you will submit your project here at each phase:

<https://forms.gle/CQsSEGM3GeCrL43c8>

Submit the following via the form for each phase (You can submit a phase before the due date):

1. Public GitHub Repo Link
2. Published App Link for Vercel.
3. Include a demo video link (must be under 90 seconds). Judges will only watch the first 90 seconds. You can [use NotebookLM](#) or record your demo.
4. WhatsApp number (top submissions will be invited to present live)

Everyone is welcome to join the Zoom meeting to watch the presentations. Only invited participants will present their submissions. The meetings start at 8:00 PM on Sundays.

Join Zoom Meeting

- Time: 08:00 PM On Sundays, December 7, 14, and 21, 2025 and on January 4, 2026 starting at 8:00 PM on Zoom. Final Live Presentation date to be determined.
- https://us06web.zoom.us/j/84976847088?pwd=Z7t7NaeXwVmmR5fysCv7NiMbfbhld_a.1
- Meeting ID: 849 7684 7088
- Passcode: 305850

Project Details: The Evolution of Todo

Focus and Theme: From CLI to Distributed Cloud-Native AI Systems.

Goal: Students act as Product Architects, using AI to build progressively complex software without writing boilerplate code.

Project Overview

This project simulates the real-world evolution of software. You will start with a simple script and end with a Kubernetes-managed, event-driven, AI-powered distributed system.

Phase Breakdown

Phase I: Todo In-Memory Python Console App

Basic Level Functionality

Objective: Build a command-line todo application that stores tasks in memory using Claude Code and Spec-Kit Plus.

Requirements

- Implement all 5 Basic Level features (Add, Delete, Update, View, Mark Complete)
- Use spec-driven development with Claude Code and Spec-Kit Plus
- Follow clean code principles and proper Python project structure

Technology Stack

- UV
- Python 3.13+
- Claude Code
- GitHub Spec-Kit

Deliverables

1. GitHub repository with:
 - Constitution file
 - specs history folder containing all specification files
 - /src folder with Python source code
 - README.md with setup instructions
 - CLAUDE.md with Claude Code instructions
2. Working console application demonstrating:
 - Adding tasks with title and description
 - Listing all tasks with status indicators
 - Updating task details
 - Deleting tasks by ID
 - Marking tasks as complete/incomplete

Windows Users: WSL 2 Setup

Windows users must use WSL 2 (Windows Subsystem for Linux) for development:

```
# Install WSL 2
wsl --install

# Set WSL 2 as default
wsl --set-default-version 2

# Install Ubuntu
wsl --install -d Ubuntu-22.04
```

Phase II: Todo Full-Stack Web Application

Basic Level Functionality

Objective: Using Claude Code and Spec-Kit Plus transform the console app into a modern multi-user web application with persistent storage.

Requirements

- Implement all 5 Basic Level features as a web application
- Create RESTful API endpoints
- Build responsive frontend interface
- Store data in Neon Serverless PostgreSQL database
- Authentication – Implement user signup/signin using Better Auth

Technology Stack

Layer	Technology
Frontend	Next.js 16+ (App Router)
Backend	Python FastAPI
ORM	SQLModel
Database	Neon Serverless PostgreSQL
Spec-Driven	Claude Code + Spec-Kit Plus
Authentication	Better Auth

API Endpoints

Method	Endpoint	Description
GET	/api/{user_id}/tasks	List all tasks
POST	/api/{user_id}/tasks	Create a new task
GET	/api/{user_id}/tasks/{id}	Get task details
PUT	/api/{user_id}/tasks/{id}	Update a task
DELETE	/api/{user_id}/tasks/{id}	Delete a task
PATCH	/api/{user_id}/tasks/{id}/complete	Toggle completion

Securing the REST API

Better Auth + FastAPI Integration

The Challenge

Better Auth is a JavaScript/TypeScript authentication library that runs on your **Next.js frontend**. However, your **FastAPI backend** is a separate Python service that needs to verify which user is making API requests.

The Solution: JWT Tokens

Better Auth can be configured to issue **JWT (JSON Web Token)** tokens when users log in. These tokens are self-contained credentials that include user information and can be verified by any service that knows the secret key.

How It Works

- User logs in on Frontend → Better Auth creates a session and issues a JWT token
- Frontend makes API call → Includes the JWT token in the Authorization: Bearer <token> header
- Backend receives request → Extracts token from header, verifies signature using shared secret
- Backend identifies user → Decodes token to get user ID, email, etc. and matches it with the user ID in the URL
- Backend filters data → Returns only tasks belonging to that user

What Needs to Change

Component	Changes Required
Better Auth Config	Enable JWT plugin to issue tokens
Frontend API Client	Attach JWT token to every API request header
FastAPI Backend	Add middleware to verify JWT and extract user
API Routes	Filter all queries by the authenticated user's ID

The Shared Secret

Both frontend (Better Auth) and backend (FastAPI) must use the **same secret key** for JWT signing and verification. This is typically set via environment variable **BETTER_AUTH_SECRET** in both services.

Security Benefits

Benefit	Description
User Isolation	Each user only sees their own tasks
Stateless Auth	Backend doesn't need to call frontend to verify users
Token Expiry	JWTs expire automatically (e.g., after 7 days)
No Shared DB Session	Frontend and backend can verify auth independently

API Behavior Change

After Auth:

All endpoints require valid JWT token
Requests without token receive 401 Unauthorized
Each user only sees/modifies their own tasks
Task ownership is enforced on every operation

Bottom Line

The REST API endpoints stay the same (**GET /api/user_id/tasks**, **POST /api/user_id/tasks**, etc.), but every request now must include a JWT token, and all responses are filtered to only include that user's data.

Monorepo Organization For Full-Stack Projects With GitHub Spec-Kit + Claude Code

This guide explains how to organize your Full-Stack Projects in a monorepo to integrate **GitHub Spec-Kit** for spec-driven development with **Claude Code**. This guide explains how to organize your repository so that Claude Code and Spec-Kit Plus can effectively edit both frontend (Next.js) and backend (FastAPI) code in a single context.

Spec-Kit Monorepo Folder Structure

```

hackathon-todo/
  -- .spec-kit/
    -- config.yaml
  -- specs/
    -- overview.md
    -- architecture.md
    -- features/
      -- task-crud.md
      -- authentication.md
      -- chatbot.md
    -- api/
      -- rest-endpoints.md
      -- mcp-tools.md
    -- database/
      -- schema.md
    -- ui/
      -- components.md
      -- pages.md
  -- CLAUXE.md
  -- frontend/
    -- CLAUXE.md
    ... (Next.js app)
  -- backend/
    -- CLAUXE.md
    ... (FastAPI app)
  -- docker-compose.yml
  -- README.md

```

Key Differences from Basic Monorepo

Aspect	Without Spec-Kit	With Spec-Kit
Specs Location	/specs (flat)	/specs (organized by type)
Config File	None	./.spec-kit/config.yaml
Spec Format	Freeform markdown	Spec-Kit conventions
Referencing	@specs/file.md	@specs/features/file.md

Spec-Kit Config File

```

# .spec-kit/config.yaml
name: hackathon-todo
version: "1.0"

structure:
  specs_dir: specs
  features_dir: specs/features
  api_dir: specs/api
  database_dir: specs/database
  ui_dir: specs/ui

phases:

```

```
- name: phase1-console
  features: [task-crud]
- name: phase2-web
  features: [task-crud, authentication]
- name: phase3-chatbot
  features: [task-crud, authentication, chatbot]
```

CLAUDE.md Files

Create multiple CLAUDE.md files to provide context at different levels:

Root CLAUDE.md

```
# Todo App - Hackathon II

## Project Overview
This is a monorepo using GitHub Spec-Kit for spec-driven development.

## Spec-Kit Structure
Specifications are organized in /specs:
- /specs/overview.md - Project overview
- /specs/features/ - Feature specs (what to build)
- /specs/api/ - API endpoint and MCP tool specs
- /specs/database/ - Schema and model specs
- /specs/ui/ - Component and page specs

## How to Use Specs
1. Always read relevant spec before implementing
2. Reference specs with: @specs/features/task-crud.md
3. Update specs if requirements change

## Project Structure
- /frontend - Next.js 14 app
- /backend - Python FastAPI server

## Development Workflow
1. Read spec: @specs/features/[feature].md
2. Implement backend: @backend/CLAUDE.md
3. Implement frontend: @frontend/CLAUDE.md
4. Test and iterate

## Commands
- Frontend: cd frontend && npm run dev
- Backend: cd backend && uvicorn main:app --reload
- Both: docker-compose up
```

Frontend CLAUDE.md

```
# Frontend Guidelines

## Stack
- Next.js 14 (App Router)
- TypeScript
- Tailwind CSS

## Patterns
- Use server components by default
- Client components only when needed (interactivity)
- API calls go through `/lib/api.ts`

## Component Structure
- `components` - Reusable UI components
- `app` - Pages and layouts

## API Client
All backend calls should use the api client:

import { api } from '@/lib/api'
const tasks = await api.getTasks()

## Styling
- Use Tailwind CSS classes
```

- No inline styles
- Follow existing component patterns

Backend CLAUDE.md

```
# Backend Guidelines

## Stack
- FastAPI
- SQLAlchemy (ORM)
- Neon PostgreSQL

## Project Structure
- `main.py` - FastAPI app entry point
- `models.py` - SQLAlchemy database models
- `routes/` - API route handlers
- `db.py` - Database connection

## API Conventions
- All routes under `/api/`
- Return JSON responses
- Use Pydantic models for request/response
- Handle errors with HTTPException

## Database
- Use SQLAlchemy for all database operations
- Connection string from environment variable: DATABASE_URL

## Running
uvicorn main:app --reload --port 8000
```

Example Spec Files

/specs/overview.md

```
# Todo App Overview

## Purpose
A todo application that evolves from console app to AI chatbot.

## Current Phase
Phase II: Full-Stack Web Application

## Tech Stack
- Frontend: Next.js 14, TypeScript, Tailwind CSS
- Backend: FastAPI, SQLAlchemy, Neon PostgreSQL
- Auth: Better Auth with JWT

## Features
- [ ] Task CRUD operations
- [ ] User authentication
- [ ] Task filtering and sorting
```

/specs/features/task-crud.md

```
# Feature: Task CRUD Operations
```

```
## User Stories
- As a user, I can create a new task
- As a user, I can view all my tasks
- As a user, I can update a task
- As a user, I can delete a task
- As a user, I can mark a task complete

## Acceptance Criteria

### Create Task
- Title is required (1-200 characters)
- Description is optional (max 1000 characters)
- Task is associated with logged-in user

### View Tasks
- Only show tasks for current user
- Display title, status, created date
- Support filtering by status
```

/specs/api/rest-endpoints.md

```
# REST API Endpoints

## Base URL
- Development: http://localhost:8000
- Production: https://api.example.com

## Authentication
All endpoints require JWT token in header:
Authorization: Bearer <token>

## Endpoints

### GET /api/tasks
List all tasks for authenticated user.

Query Parameters:
- status: "all" | "pending" | "completed"
- sort: "created" | "title" | "due_date"

Response: Array of Task objects

### POST /api/tasks
Create a new task.

Request Body:
- title: string (required)
- description: string (optional)

Response: Created Task object
```

/specs/database/schema.md

```
# Database Schema

## Tables

### users (managed by Better Auth)
- id: string (primary key)
- email: string (unique)
- name: string
- created_at: timestamp

### tasks
- id: integer (primary key)
- user_id: string (foreign key -> users.id)
- title: string (not null)
- description: text (nullable)
- completed: boolean (default false)
- created_at: timestamp
- updated_at: timestamp

## Indexes
- tasks.user_id (for filtering by user)
- tasks.completed (for status filtering)
```

Workflow with Spec-Kit + Claude Code

- Write/Update Spec → @specs/features/new-feature.md
- Ask Claude Code to Implement → "Implement @specs/features/new-feature.md"
- Claude Code reads: Root CLAUDE.md, Feature spec, API spec, Database spec, Relevant CLAUDE.md
- Claude Code implements in both frontend and backend
- Test and iterate on spec if needed

Referencing Specs in Claude Code

```
# Implement a feature
You: @specs/features/task-crud.md implement the create task feature

# Implement API
You: @specs/api/rest-endpoints.md implement the GET /api/tasks endpoint

# Update database
You: @specs/database/schema.md add due_date field to tasks

# Full feature across stack
You: @specs/features/authentication.md implement Better Auth login
```

Summary

Component	Purpose
<code>/.spec-kit/config.yaml</code>	Spec-Kit configuration
<code>/specs/overview.md</code>	Project overview and status
<code>/specs/features/</code>	What to build (user stories, acceptance criteria)
<code>/specs/api/</code>	How APIs should work
<code>/specs/database/</code>	Data models and schema
<code>/specs/ui/</code>	UI components and pages
<code>/CLAUDE.md</code>	How to navigate and use specs
<code>/frontend/CLAUDE.md</code>	Frontend-specific patterns
<code>/backend/CLAUDE.md</code>	Backend-specific patterns

Key Point:

Spec-Kit provides organized, structured specs that Claude Code can reference. The CLAUDE.md files tell Claude Code how to use those specs and project-specific conventions.

Summary: Monorepo vs Separate Repos

Approach	Pros	Cons
Monorepo ★	Single CLAUDE.md context, easier cross-cutting changes	Larger repo
Separate Repos	Clear separation, independent deployments	Claude Code needs workspace setup

Recommendation:

Use monorepo for the hackathon – simpler for Claude Code to navigate and edit both frontend and backend in a single context.

Key Benefits of This Structure

Benefit	Description
Single Context	Claude Code sees entire project, can make cross-cutting changes
Layered CLAUDE.md	Root file for overview, subfolder files for specific guidelines
Specs Folder	Reference specifications directly with @specs/filename.md
Clear Separation	Frontend and backend code in separate folders, easy to navigate

Phase III: Todo AI Chatbot

Basic Level Functionality

Objective: Create an AI-powered chatbot interface for managing todos through natural language using MCP (Model Context Protocol) server architecture and using Claude Code and Spec-Kit Plus.


Requirements


1. Implement conversational interface for all Basic Level features
2. Use OpenAI Agents SDK for AI logic
3. Build MCP server with Official MCP SDK that exposes task operations as tools
4. Stateless chat endpoint that persists conversation state to database
5. AI agents use MCP tools to manage tasks. The MCP tools will also be stateless and will store state in the database.

Technology Stack

Component	Technology
Frontend	OpenAI ChatKit
Backend	Python FastAPI
AI Framework	OpenAI Agents SDK
MCP Server	Official MCP SDK
ORM	SQLModel
Database	Neon Serverless PostgreSQL
Authentication	Better Auth

Architecture





Database Models

Model	Fields	Description
Task	user_id, id, title, description, completed, created_at, updated_at	Todo items
Conversation	user_id, id, created_at, updated_at	Chat session
Message	user_id, id, conversation_id, role (user/assistant), content, created_at	Chat history

Chat API Endpoint

Method	Endpoint	Description
POST	/api/{user_id}/chat	Send message & get AI response

Request

Field	Type	Required	Description
conversation_id	integer	No	Existing conversation ID (creates new if not provided)
message	string	Yes	User's natural language message

Response

Field	Type	Description
conversation_id	integer	The conversation ID
response	string	AI assistant's response
tool_calls	array	List of MCP tools invoked

MCP Tools Specification

The MCP server must expose the following tools for the AI agent:

Tool: add_task

Purpose	Create a new task
Parameters	user_id (string, required), title (string, required), description (string, optional)
Returns	task_id, status, title
Example Input	{"user_id": "ziakhan", "title": "Buy groceries", "description": "Milk, eggs, bread"}

Example Output	{"task_id": 5, "status": "created", "title": "Buy groceries"}
-----------------------	---

Tool: list_tasks

Purpose	Retrieve tasks from the list
Parameters	status (string, optional: "all", "pending", "completed")
Returns	Array of task objects
Example Input	{"user_id": (string, required), "status": "pending"}
Example Output	[{"id": 1, "title": "Buy groceries", "completed": false}, ...]

Tool: complete_task

Purpose	Mark a task as complete
Parameters	user_id (string, required), task_id (integer, required)
Returns	task_id, status, title
Example Input	{"user_id": "ziakhan", "task_id": 3}
Example Output	{"task_id": 3, "status": "completed", "title": "Call mom"}

Tool: delete_task

Purpose	Remove a task from the list
Parameters	user_id (string, required), task_id (integer, required)
Returns	task_id, status, title
Example Input	{"user_id": "ziakhan", "task_id": 2}
Example Output	{"task_id": 2, "status": "deleted", "title": "Old task"}

Tool: update_task

Purpose	Modify task title or description
Parameters	user_id (string, required), task_id (integer, required), title (string, optional), description (string, optional)
Returns	task_id, status, title
Example Input	{"user_id": "ziakhan", "task_id": 1, "title": "Buy groceries and fruits"}
Example Output	{"task_id": 1, "status": "updated", "title": "Buy groceries and fruits"}

Agent Behavior Specification

Behavior	Description
Task Creation	When user mentions adding/creating/remembering something, use add_task
Task Listing	When user asks to see/show/list tasks, use list_tasks with appropriate filter
Task Completion	When user says done/complete/finished, use complete_task
Task Deletion	When user says delete/remove/cancel, use delete_task
Task Update	When user says change/update/rename, use update_task
Confirmation	Always confirm actions with friendly response
Error Handling	Gracefully handle task not found and other errors

Conversation Flow (Stateless Request Cycle)

1. Receive user message
2. Fetch conversation history from database
3. Build message array for agent (history + new message)
4. Store user message in database
5. Run agent with MCP tools
6. Agent invokes appropriate MCP tool(s)
7. Store assistant response in database
8. Return response to client
9. Server holds NO state (ready for next request)

Natural Language Commands

The chatbot should understand and respond to:

User Says	Agent Should
"Add a task to buy groceries"	Call add_task with title "Buy groceries"
"Show me all my tasks"	Call list_tasks with status "all"
"What's pending?"	Call list_tasks with status "pending"
"Mark task 3 as complete"	Call complete_task with task_id 3
"Delete the meeting task"	Call list_tasks first, then delete_task
"Change task 1 to 'Call mom tonight'"	Call update_task with new title
"I need to remember to pay bills"	Call add_task with title "Pay bills"
"What have I completed?"	Call list_tasks with status "completed"

Deliverables

1. GitHub repository with:
 - /frontend – ChatKit-based UI
 - /backend – FastAPI + Agents SDK + MCP
 - /specs – Specification files for agent and MCP tools
 - Database migration scripts
 - README with setup instructions
2. Working chatbot that can:
 - Manage tasks through natural language via MCP tools
 - Maintain conversation context via database (stateless server)
 - Provide helpful responses with action confirmations
 - Handle errors gracefully
 - Resume conversations after server restart

Key Architecture Benefits

Aspect	Benefit
MCP Tools	Standardized interface for AI to interact with your app
Single Endpoint	Simpler API — AI handles routing to tools
Stateless Server	Scalable, resilient, horizontally scalable
Tool Composition	Agent can chain multiple tools in one turn

Key Stateless Architecture Benefits

- **Scalability:** Any server instance can handle any request
- **Resilience:** Server restarts don't lose conversation state
- **Horizontal scaling:** Load balancer can route to any backend
- **Testability:** Each request is independent and reproducible

Phase IV: Local Kubernetes Deployment (Minikube, Helm Charts, kubectl-ai, Kagent, Docker Desktop, and Gordon)

Cloud Native Todo Chatbot with Basic Level Functionality

Objective: Deploy the Todo Chatbot on a local Kubernetes cluster using Minikube, Helm Charts.

Requirements

- Containerize frontend and backend applications (Use Gordon)
- Use Docker AI Agent (Gordon) for AI-assisted Docker operations
- Create Helm charts for deployment (Use kubectl-ai and/or kagent to generate)
- Use kubectl-ai and kagent for AI-assisted Kubernetes operations
- Deploy on Minikube locally

Note: If Docker AI (Gordon) is unavailable in your region or tier, use standard Docker CLI commands or ask Claude Code to generate the `docker run` commands for you.

Technology Stack

Component	Technology
Containerization	Docker (Docker Desktop)
Docker AI	Docker AI Agent (Gordon)
Orchestration	Kubernetes (Minikube)
Package Manager	Helm Charts
AI DevOps	kubectl-ai, and Kagent
Application	Phase III Todo Chatbot

AIOps

Use [Docker AI Agent \(Gordon\)](#) for intelligent Docker operations:

```
# To know its capabilities
docker ai "What can you do?"
```

Enable Gordon: Install latest Docker Desktop 4.53+, go to Settings > Beta features, and toggle it on.

Use [kubectl-ai](#), and [Kagent](#) for intelligent Kubernetes operations:

```
# Using kubectl-ai
kubectl-ai "deploy the todo frontend with 2 replicas"
kubectl-ai "scale the backend to handle more load"
kubectl-ai "check why the pods are failing"

# Using kagent
kagent "analyze the cluster health"
kagent "optimize resource allocation"
```

Starting with kubectl-ai will make you feel empowered from day one. Layer in Kagent for advanced use cases. Pair them with Minikube for zero-cost learning and work.

Research Note: Using Blueprints for Spec-Driven Deployment

Can Spec-Driven Development be used for infrastructure automation, and how we may need to use blueprints powered by Claude Code Agent Skills.

1. [Is Spec-Driven Development Key for Infrastructure Automation?](#)
2. [ChatGPT Progressive Learning Conversation](#)
3. [Spec-Driven Cloud-Native Architecture: Governing AI Agents for Managed Services with Claude Code and SpecKit](#)

Phase V: Advanced Cloud Deployment

Advanced Level Functionality on DigitalOcean Kubernetes

Objective: Implement advanced features and deploy first on Minikube locally and then to production-grade Kubernetes on DigitalOcean and Kafka on Redpanda Cloud.

Part A: Advanced Features

- Implement all Advanced Level features (Recurring Tasks, Due Dates & Reminders)
- Implement Intermediate Level features (Priorities, Tags, Search, Filter, Sort)
- Add event-driven architecture with Kafka
- Implement Dapr for distributed application runtime

Part B: Local Deployment

- Deploy to Minikube
- Deploy Dapr on Minikube use Full Dapr: Pub/Sub, State, Bindings (cron), Secrets, Service Invocation

Part C: Cloud Deployment

- Deploy to DigitalOcean Kubernetes (DOKS)
- Deploy Dapr on DOKS use Full Dapr: Pub/Sub, State, Bindings (cron), Secrets, Service Invocation
- Use Kafka on Redpanda Cloud
- Set up CI/CD pipeline using Github Actions
- Configure monitoring and logging

DigitalOcean Setup


New DigitalOcean accounts receive \$200 credit for 60 days:

1. Sign up at digitalocean.com
2. Create a Kubernetes cluster (DOKS)
3. Configure kubectl to connect to DOKS
4. Deploy using Helm charts from Phase IV

Kafka Use Cases in Phase


Event-Driven Architecture for Todo Chatbot

1. Reminder/Notification System




When a task with a due date is created, publish a reminder event. A separate notification service consumes and sends reminders at the right time.

2. Recurring Task Engine




When a recurring task is marked complete, publish an event. A separate service consumes it and auto-creates the next occurrence.

3. Activity/Audit Log




Every task operation (create, update, delete, complete) publishes to Kafka. An audit service consumes and maintains a complete history.


4. Real-time Sync Across Clients



Changes from one client are broadcast to all connected clients in real-time.

Recommended Architecture





Kafka Topics

Topic	Producer	Consumer	Purpose
task-events	Chat API (MCP Tools)	Recurring Task Service, Audit Service	All task CRUD operations
reminders	Chat API (when due date set)	Notification Service	Scheduled reminder triggers
task-updates	Chat API	WebSocket Service	Real-time client sync

Event Schema Examples

Task Event

Field	Type	Description
event_type	string	"created", "updated", "completed", "deleted"
task_id	integer	The task ID
task_data	object	Full task object
user_id	string	User who performed action
timestamp	datetime	When event occurred

Reminder Event

Field	Type	Description
task_id	integer	The task ID
title	string	Task title for notification
due_at	datetime	When task is due
remind_at	datetime	When to send reminder
user_id	string	User to notify

Why Kafka for Todo App?

Without Kafka	With Kafka
Reminder logic coupled with main app	Decoupled notification service
Recurring tasks processed synchronously	Async processing, no blocking
No activity history	Complete audit trail
Single client updates	Real-time multi-client sync
Tight coupling between services	Loose coupling, scalable

Bottom Line

Kafka turns the Todo app from a simple CRUD app into an **event-driven system** where services communicate through events rather than direct API calls. This is essential for the advanced features (recurring tasks, reminders) and scales better in production.

Key Takeaway:

Kafka enables decoupled, scalable microservices architecture where the Chat API publishes events and specialized services (Notification, Recurring Task, Audit) consume and process them independently.

Kafka Service Recommendations

For Cloud Deployment

Service	Free Tier	Pros	Cons
Redpanda Cloud 	Free Serverless tier	Kafka-compatible, no Zookeeper, fast, easy setup	Newer ecosystem
Confluent Cloud	\$400 credit for 30 days	Industry standard, Schema Registry, great docs	Credit expires
CloudKarafka	"Developer Duck" free plan	Simple, 5 topics free	Limited throughput
Aiven	\$300 credit trial	Fully managed, multi-cloud	Trial expires
Self-hosted (Strimzi)	Free (just compute cost)	Full control, learning experience	More complex setup

For Local Development (Minikube)

Option	Complexity	Description
Redpanda (Docker) 	Easy	Single binary, no Zookeeper, Kafka-compatible
Bitnami Kafka Helm	Medium	Kubernetes-native, Helm chart
Strimzi Operator	Medium-Hard	Production-grade K8s operator

Primary Recommendation: Redpanda Cloud (Serverless)

Best for hackathon because:

- Free serverless tier (no credit card for basic usage)
- Kafka-compatible - same APIs, clients work unchanged
- No Zookeeper - simpler architecture
- Fast setup - under 5 minutes
- REST API + Native protocols

Sign up: <https://redpanda.com/cloud>

For Local/Minikube: Redpanda Docker

Single container, Kafka-compatible:

```
# docker-compose.redpanda.yml
services:
  redpanda:
    image: redpandadata/redpanda:latest
    command:
      - redpanda start
      - --smp 1
      - --memory 512M
      - --overprovisioned
      - --kafka-addr PLAINTEXT://0.0.0.0:9092
      - --advertise-kafka-addr PLAINTEXT://localhost:9092
    ports:
      - "9092:9092"
      - "8081:8081" # Schema Registry
      - "8082:8082" # REST Proxy
```

Alternative: Self-Hosted on Kubernetes (Strimzi)

Good learning experience for students:

```
# Install Strimzi operator
kubectl create namespace kafka
kubectl apply -f https://strimzi.io/install/latest?namespace=kafka

# Create Kafka cluster
kubectl apply -f kafka-cluster.yaml
```

Redpanda Cloud Quick Setup

Step	Action
1	Sign up at redpanda.com/cloud
2	Create a Serverless cluster (free tier)
3	Create topics: task-events, reminders, task-updates
4	Copy bootstrap server URL and credentials
5	Use standard Kafka clients (kafka-python, aiokafka)

Python Client Example

Standard kafka-python works with Redpanda:

```
from kafka import KafkaProducer
import json

producer = KafkaProducer(
    bootstrap_servers="YOUR-CLUSTER.cloud.redpanda.com:9092",
    security_protocol="SASL_SSL",
    sasl_mechanism="SCRAM-SHA-256",
    sasl_plain_username="YOUR-USERNAME",
    sasl_plain_password="YOUR-PASSWORD",
    value_serializer=lambda v: json.dumps(v).encode('utf-8')
)

# Publish event
producer.send("task-events", {"event_type": "created", "task_id": 1})
```

Summary for Hackathon

Type	Recommendation
Local: Minikube	Redpanda Docker container

Type	Recommendation
Cloud	Redpanda Cloud Serverless (free) or Strimzi self-hosted

Dapr Integration Guide

What is Dapr?


Dapr (Distributed Application Runtime) is a portable, event-driven runtime that simplifies building microservices. It runs as a **sidecar** next to your application and provides building blocks via HTTP/gRPC APIs.

Dapr Building Blocks for Todo App


Building Block	Use Case in Todo App
Pub/Sub	Kafka abstraction – publish/subscribe without Kafka client code
State Management	Conversation state storage (alternative to direct DB calls)
Service Invocation	Frontend → Backend communication with built-in retries
Bindings	Cron triggers for scheduled reminders
Secrets Management	Store API keys, DB credentials securely

Architecture: Without Dapr vs With Dapr

Without Dapr (Direct Dependencies)



With Dapr (Abstracted Dependencies)



Use Case 1: Pub/Sub (Kafka Abstraction)

Instead of using kafka-python directly, publish events via Dapr:

Without Dapr:

```
from kafka import KafkaProducer
producer = KafkaProducer(bootstrap_servers="kafka:9092", ...)
producer.send("task-events", value=event)
```

With Dapr:

```
import httpx

# Publish via Dapr sidecar (no Kafka library needed!)
await httpx.post(
    "http://localhost:3500/v1.0/publish/kafka-pubsub/task-events",
    json={"event_type": "created", "task_id": 1}
)
```

Dapr Component Configuration:

```
apiVersion: dapr.io/v1alpha1
kind: Component
metadata:
  name: kafka-pubsub
spec:
  type: pubsub.kafka
  version: v1
  metadata:
    - name: brokers
      value: "kafka:9092"
    - name: consumerGroup
      value: "todo-service"
```

Use Case 2: State Management (Conversation State)

Store conversation history without direct DB code:

Without Dapr:

```
from sqlmodel import Session
session.add(Message(...))
session.commit()
```

With Dapr:

```
import httpx

# Save state via Dapr
await httpx.post(
    "http://localhost:3500/v1.0/state/statesstore",
    json=[{
        "key": f"conversation-{conv_id}",
        "value": {"messages": messages}
    }]
)

# Get state
response = await httpx.get(
    f"http://localhost:3500/v1.0/state/statesstore/conversation-{conv_id}"
)
```

Dapr Component Configuration:

```
apiVersion: dapr.io/v1alpha1
kind: Component
metadata:
  name: statesstore
spec:
  type: state.postgresql
```

```
version: v1
metadata:
  - name: connectionString
    value: "host=neon.db user=... password=... dbname=todo"
```

Use Case 3: Service Invocation (Frontend → Backend)

Built-in service discovery, retries, and mTLS:

Without Dapr:

```
// Frontend must know backend URL
fetch("http://backend-service:8000/api/chat", {...})
```

With Dapr:

```
// Frontend calls via Dapr sidecar - automatic discovery
fetch("http://localhost:3500/v1.0/invoke/backend-service/method/api/chat",
{...})
```

Use Case 4: Input Bindings (Scheduled Reminders)

Trigger reminder checks on a schedule:

Dapr Cron Binding:

```
apiVersion: dapr.io/v1alpha1
kind: Component
metadata:
  name: reminder-cron
spec:
  type: bindings.cron
  version: v1
  metadata:
    - name: schedule
      value: "*/5 * * * *" # Every 5 minutes
```

Backend Handler:

```
@app.post("/reminder-cron")
async def check_reminders():
    # Dapr calls this every 5 minutes
    # Check for due tasks and send notifications
    pass
```

Use Case 5: Secrets Management

Securely store and access credentials:

Dapr Component (Kubernetes Secrets):


```
apiVersion: dapr.io/v1alpha1
kind: Component
metadata:
  name: kubernetes-secrets
spec:
  type: secretstores.kubernetes
  version: v1
```

Access in App:

```
import httpx

response = await httpx.get(
    "http://localhost:3500/v1.0/secrets/kubernetes-secrets/openai-api-key"
)
api_key = response.json()["openai-api-key"]
```

Complete Dapr Architecture



Dapr Components Summary

Component	Type	Purpose
kafka-pubsub	pubsub.kafka	Event streaming (task-events, reminders)
statestore	state.postgresql	Conversation state, task cache
reminder-cron	bindings.cron	Trigger reminder checks
kubernetes-secrets	secretstores.kubernetes	API keys, DB credentials

Why Use Dapr?

Without Dapr	With Dapr
Import Kafka, Redis, Postgres libraries	Single HTTP API for all
Connection strings in code	Dapr components (YAML config)
Manual retry logic	Built-in retries, circuit breakers
Service URLs hardcoded	Automatic service discovery
Secrets in env vars	Secure secret store integration
Vendor lock-in	Swap Kafka for RabbitMQ with config change

Local vs Cloud Dapr Usage

Phase	Dapr Usage
Local (Minikube)	Install Dapr, use Pub/Sub for Kafka, basic state management
Cloud (DigitalOcean)	Full Dapr: Pub/Sub, State, Bindings (cron), Secrets, Service Invocation

Getting Started with Dapr

```
# Install Dapr CLI
curl -fsSL https://raw.githubusercontent.com/dapr/cli/master/install/install.sh
| bash

# Initialize Dapr on Kubernetes
dapr init -k

# Deploy components
kubectl apply -f dapr-components/

# Run app with Dapr sidecar
dapr run --app-id backend --app-port 8000 -- uvicorn main:app
```

Bottom Line

Dapr abstracts infrastructure (Kafka, DB, Secrets) behind simple HTTP APIs. Your app code stays clean, and you can swap backends (e.g., Kafka → RabbitMQ) by changing YAML config, not code.

Submission Requirements

Required Submissions

1. Public GitHub Repository containing:
 - All source code for all completed phases
 - /specs folder with all specification files
 - CLAUDE.md with Claude Code instructions
 - README.md with comprehensive documentation
 - Clear folder structure for each phase
2. Deployed Application Links:
 - Phase II: Vercel/frontend URL + Backend API URL
 - Phase III-V: Chatbot URL
 - Phase IV: Instructions for local Minikube setup
 - Phase V: DigitalOcean deployment URL
3. Demo Video (maximum 90 seconds):
 - Demonstrate all implemented features
 - Show spec-driven development workflow
 - Judges will only watch the first 90 seconds
4. WhatsApp Number for presentation invitation

Resources

Core Tools

Tool	Link	Description
Claude Code	claude.com/product/clause-code	AI coding assistant
GitHub Spec-Kit	github.com/panaversity/spec-kit-plus	Specification management
OpenAI ChatKit	platform.openai.com/docs/guides/chat-kit	Chatbot UI framework
MCP	github.com/modelcontextprotocol/python-sdk	MCP server framework

Infrastructure

Service	Link	Notes
Neon DB	neon.tech	Free tier available
Vercel	vercel.com	Free frontend hosting
DigitalOcean	digitalocean.com	\$200 credit for 60 days
Minikube	minikube.sigs.k8s.io	Local Kubernetes

Frequently Asked Questions

Q: Can I skip phases?

A: No, each phase builds on the previous. You must complete them in order.

Q: Can I use different technologies?

A: The core stack must remain as specified. You can add additional tools/libraries.

Q: Do I need a DigitalOcean account from the start?

A: No, only for Phase V. Use the \$200 free credit for new accounts.

Q: Can I work in a team?

A: This is an individual hackathon. Each participant submits separately.

Q: What if I don't complete all the phases?

A: Submit what you complete. Partial submissions are evaluated proportionally.

Good luck, and may your specs be clear and your code be clean!



— *The Panaversity, PIAIC, and GIAIC Teams*